Assignment 3 - Reinforcement Learning

Michał Jóźwik, Lukas Ostrovskis, Martynas Krupskis, Denis Tsvetkov

March 2022

1 Running the code

Open **Group_75_code**/**main**/**mysolution** directory with PyCharm, inside you will find RunMe.py file, which you can run through PyCharm. You can also run this code via command line using *python RunMe.py* command (tested with Python 3.9.5) in the provided directory. You need to install numpy before running the code. The code in RunMe.py will find the path in the *toy_maze.txt* and print it's length to the console.

2 Development

1. To ensure that the agent is not biased towards selecting the same action in **getBestAction()**, we first check whether all the action values within the valid action list are equal. If they are, we select the action randomly, otherwise we choose the best one.

2. We train the agent for a defined number of steps (30000 for the toy maze and 130000 for the easy one). The agent starts taking steps until it reaches its goal and then its position is reset and the number of steps left that it can make is reduced. For each step, we get the current state of the agent, use the e-greedy algorithm to select the next action, perform that action to obtain the next state, and finally use these obtained values to update the Q-table for the old state.

3. It stops after 30000 steps.

4. As it can be seen in the plots below the average number of steps to reach a target per trial randomly fluctuates between 100 and 2500 for the toy maze and 4000 and 11000 for the easy one but with no clear pattern of downtrend. Thus, we can conclude that the agent is in fact not learning anything yet.

Figure 1: Average number of steps per trial for the toy maze for 30000 steps

Figure 2: Average number of steps per trial for the easy maze for 130000 steps

5. This is how our **updateQ()** method works:

- Get the **Q(s', a')** for each possible action **a'** from next state **s'** and store that into **action_values** variable.
- Get the maximum action value from next state **s'** using **action_values[argmax(action_values]** and store it to **Q_max**.
- Get the **Q** value of the current state **s** and the action **a** and store it to **Q_old**

2

- Compute $\mathbf{Q\_new} = \mathbf{Q\_old} + \boldsymbol{\alpha} \cdot (\boldsymbol{r} + \boldsymbol{\gamma} \cdot \mathbf{Q\_max} - \mathbf{Q\_old})$, formula given to us in the lecture and in the assignment itself.

- Set $\mathbf{Q(s, a)}$ to $\mathbf{Q\_new}$ computed in step 4.

6. In figure 3 and 4 we can see that our agent is now able to learn both mazes, as there is a clear downtrend for the average number of steps (it sometimes goes up but that is due to the randomness of $\boldsymbol{\epsilon}$). Furthermore, both graphs are converging to an optimal value in around 35 trials for the toy maze and 45 for the easy one.


Figure 3: Average number of steps per trial for the toy maze for 30000 steps[1]

---

[1]We introduced an additional stopping criterion of 100 trials (only for the plotting but the not the final version of the algorithm) so that the learning and convergence of the algorithm are clearer. This applies to all plots except figures 1, 2, 17, and 18.

Figure 4: Average number of steps per trial for the toy maze for 130000 steps

3 Training

7. We have chosen to generate plots with epsilon values of 0.1, 0.3, 0.6 and 0.9.
   Not to repeat ourselves the effects of high and low epsilon values are explained
   in question 8, meaning that going from 0.1 to 0.9 our agent is prioritizing explo-
   ration over exploitation more and more, therefore more and more random and
   non-optimal steps are likely to get included in the path. In the figures 3, 5, 6,
   7, 4, 8, 9, and 10 corresponding to epsilon values of 0.1, 0.3, 0.6, and 0.9 for
   both mazes this can be clearly seen, as the optimal path we find is getting worse
   as more and more random and non-optimal steps are introduced into our path
   from start to target location.

Figure 5: Average number of steps per trial for the toy maze for $\epsilon = \mathbf{0.3}$ and 30000 steps


Figure 6: Average number of steps per trial for the toy maze for $\epsilon = \mathbf{0.6}$ and 30000 steps

Figure 7: Average number of steps per trial for the toy maze for $\epsilon = 0.9$ and 30000 steps


Figure 8: Average number of steps per trial for the easy maze for $\epsilon = 0.3$ and 130000 steps

6

Figure 9: Average number of steps per trial for the easy maze for $\epsilon = \mathbf{0.6}$ and 130000 steps


Figure 10: Average number of steps per trial for the easy maze for $\epsilon = \mathbf{0.9}$ and 130000 steps

8. The trade-off between a high and a low epsilon is the trade-off between exploration and exploitation. The agent must exploit the actions that have given it a high reward in the past to obtain a lot of reward, but to discover such actions, it also has to try actions it has not tried before - it has to explore. High epsilon is equivalent to relying heavily on exploration, thus the agent will explore a larger solution space, but won't be as effective when learning from it. Low epsilon is equivalent to relying heavily on exploitation, thus the agent will exploit the

solutions that have already led it to rewards, but this might result in the agent getting stuck in a local optima.

9. Theoretically lower alpha values should take a larger number of trials to converge to an optimum, but in our experimentation (look at figures 11, 12, 3, 13, 14, 15, 4, 16 corresponding to the alpha values of 0.1, 0.4, 0.7, and 1 for both of the mazes) there does not seem to be any correlation between alpha values and the number of trials it takes to converge to an optimum. After some experimentation we were able to conclude that this is due to the fact that $(r + \gamma \cdot \mathbf{Q\_max} - \mathbf{Q\_old})$ in each iteration is really small, in the scale of $10e - 6$. Meaning that multiplying this value by alpha of 0.1 or 0.9 has negligible influence to the value by which the Q(s, a) is updated.

Figure 11: Average number of steps per trial for the toy maze for $\boldsymbol{\alpha = 0.1}$ and 30000 steps

Figure 12: Average number of steps per trial for the toy maze for $\alpha = 0.4$ and 30000 steps


Figure 13: Average number of steps per trial for the toy maze for $\alpha = 1$ and 30000 steps

9

Figure 14: Average number of steps per trial for the easy maze for $\alpha = 0.1$ and 130000 steps

Figure 15: Average number of steps per trial for the easy maze for $\alpha = 0.4$ and 130000 steps

10

Figure 16: Average number of steps per trial for the easy maze for $\alpha = 1$ and 130000 steps

4   Optimization

10. After adding a second reward, sized 5, our algorithm switches to converging to the coordinate with the new reward of 5, meaning that it was unable to converge to an optimal reward. This is easy to explain, as in question number 7 the optimal epsilon value is 0.1. Meaning that our algorithm currently prioritizes exploitation over exploration and it simply cannot find that (9, 9) contains a bigger reward, as finding reward of 5 at (9, 0) is easier to find.

11. Our algorithm was able to find a path to a coordinate with a reward of 10 with an epsilon value of around 0.5. Increasing the epsilon value further also increases the number of steps we have to take to find the reward of 10.
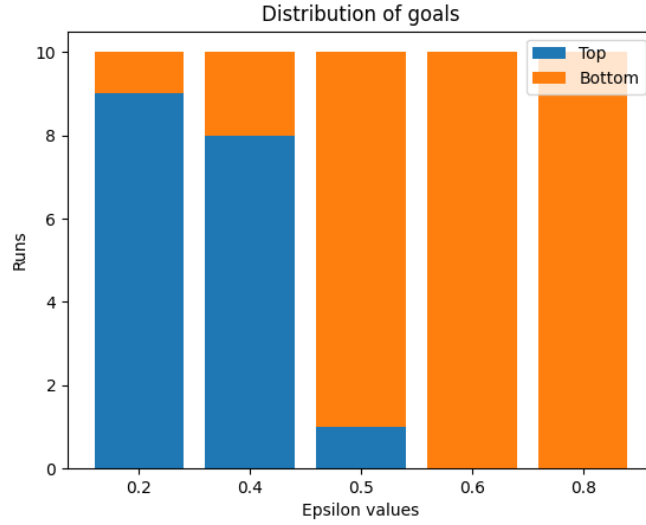
Figure 17: The distribution of the goals reached depending on the value of epsilon. Bottom is the optimal goal (the one at coordinate (9, 9) with reward 10).

12. Gamma is a parameter responsible for importance of future rewards, so 0 will make the agent myopic, meaning that it will only consider current rewards, and for gamma values approaching 1 the agent will take long-term reward into account more and more. In our maze, reaching a reward of 5 is easier, as it requires less exploration, but reaching a reward of 10 is more rewarding, therefore high gamma value means optimizing for a reward of 10. If we want to go up for a smaller reward then we have to become more short-sighted for future rewards and try to reach a target that is "easier" to reach. This is evident in the figure below, where high gamma values result in the majority of runs going down (in the maze) and reaching higher reward, but low gamma values result in going up for the smaller reward.

As it is visible in our plot, the turning point for the agent is the gamma value of around 0.75. For all gamma values below it, it tends to always end up at the top goal with the small reward, at 0.8 it starts terminating at the bottom goal more frequently, and for all values greater than or equal to 0.9 it tends to always end up at the bottom goal with the big reward.
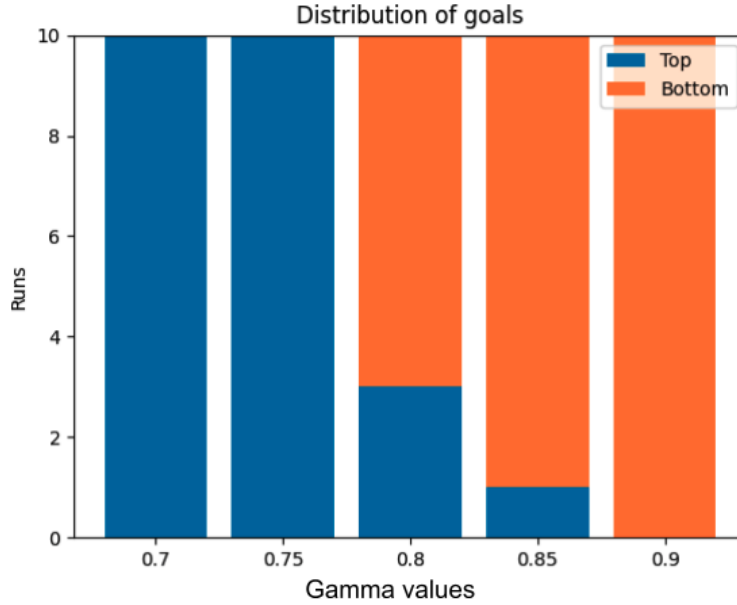
Figure 18: Top and bottom goal terminations over 10 runs for different gamma values. Bottom is the optimal goal (the one at coordinate (9, 9) with reward 10).

5 Reflection

13. The greedy actions may not always be optimal in the long run, as sometimes a slight loss in the short term may lead to the overall better solution. This could lead to being stuck in the local optimum instead of finding the global optimum that we're looking for.

14. The reward function is generally designed by humans, who have a certain goal in mind and they want the agent to get to that goal in the most efficient way. In most cases they also have a general idea of how the agent should perform, but it also happens that the system discovers unexpected ways to obtain a lot of reward. This could be considered a reward function exploitation and it could lead to unintended negative consequences, which were not thought of by the designers. As currently more companies opt for Machine Learning solutions we need to be aware that the main stakeholders of most of these projects are humans.

15. The easiest way to mitigate this problem is to carefully design the reward function and ensure that there are means to mitigate the problems that arise. Engineers have to be prepared to use these means to adjust it whenever there is an unintended solution. In general it might not always be possible to avoid all the possible exploitation techniques, but this might also lead to the discovery of better unexpected solutions. There is no silver bullet when it comes to machine learning, but there are definite ways to make it more controlled. We need to make our code readable and modifiable to possibly act whenever there are changes needed.